


DSL Embedding in Scala

Tiark Rompf

PURDUE
UNIVERSITY

DSL Summer School, July 2015



scala



Novell



SIEMENS



guardian.co.uk



Bank of America



Thatcham

Autodesk®



Scala = scalable language

Small Scripts -> Large Systems

gradual, evolutionary

FP + OOP

```
class Matrix[T:Numeric](val rows: Int, val cols: Int) {  
    private val arr: Array[T] = new Array[T](rows * cols)  
  
    def apply(i: Int, j: Int): T = arr(i*cols + j)  
  
    def *(that: Matrix[T]) = { val res = new Matrix[T]; ...; res }  
    def +(that: Matrix[T]) = ...  
}  
  
// a,b,c,d: Matrix[Double]  
val x = a*b + a*c  
val y = a*c + a*d  
println(x+y)
```

Operator overloading, higher-order functions, by-name parameters, implicits, traits, ...

What about performance?

"The compiler / JVM will
make it run fast"

(wishful thinking)

10x – 1000x slowdown

(hard reality)

News

 My Bookmarks

 Contribute an Article

 print

Yammer Moving from Scala to Java

Posted by Alex Blewitt on Nov 30, 2011

Sections [Development](#) Topics [Programming](#), [Social Networking](#)

Share 

An [e-mail](#), sent from Yammer's Typesafe, ended up being that Yammer is moving its codebase from Scala to Java with complexity and performance.

Yammer PR Shelley Risk confirmed that Coda Hale, rather than the original author has been leading the move. Coda clarified that the move was initiated by Typesafe (CEO of [Typesafe](#)) following a recent acquisition.

"Via profiling and examining the bytecode we managed to get a 100x improvement by adopting some simple rules:

Don't ever use a for-loop

Don't ever use `scala.collection.mutable`

Don't ever use `scala.collection.immutable`

Always use `private[this]`

Avoid closures"

Update: Code has published [Yammer's official position](#) on the subject, which confirms the move is intended to make the codebase more portable (switching from Scala) and that the move will be gradual.

Slow Program -> Fast Program?

Much harder!

**"Any problem in computer
science can be solved by another
level of indirection"**

-- David Wheeler

**"Any problem in computer
science can be solved by
another level of indirection --
except problems caused by too
many levels of indirection"**

Abstraction Penalty


Type Classes


```
class Matrix[T<:Numeric<:Manifest<:>], val rows: Int, val cols: Int) {  
    private val arr: Array[T] = new Array[T](rows * cols)  
    private val num = implicitly[Numeric[T]]; import num._  
  
    def apply(i: Int, j: Int): T = arr(i*cols + j)  
  
    def update(i: Int, j: Int, e: T): Unit =  
        arr(i*cols + j) = e  
  
    def *(that: Matrix[T]): Matrix[T] = {  
        val res = new Matrix[T](this.rows, that.cols)  
  
        for (i <- 0 until this.rows) {  
            for (j <- 0 until that.cols) {  
                for (k <- 0 until this.rows)  
                    res(i, j) += this(i, k) * that(k, j)  
            }  
        }  
        res  
    }  
}
```

Indirection


Closures
(and megamorphic
call sites)

Method
Calls






Principles of Compiler Design



Alfred V. Aho
Jeffrey D. Ullman

Programmer




general purpose compiler

Hardware

(illustration: Markus Püschel)

Programmer



Hardware

- Matrix, Graph, ...
- Array, Struct, Loop, ...
- SIMD, GPU, cluster, ...

- **horizontal and vertical extensibility**
- **generic optimizations** at each level (cse, dce, ...)

Secret Weapon: LMS


LMS = Lightweight Modular Staging

- An extensible compiler framework
- Implemented as a Scala library
- Execute 'now' vs 'compile and exec later'
- Specialize and compile program pieces at runtime

Staging = Multi-Stage Programming


- Computations can generally be separated into stages (Jørring, Scherlis 1986), distinguished by:
 - frequency of execution
 - availability of data
- Multi-Stage Programming (Taha, Sheard 1997): make stages explicit in a program:
 - “delay” expressions to a generated stage
 - “run” delayed expressions
 - staged program fragments as first class values

Generative **METAPROGRAMMING**



(by Amorphia Apparel)

program generically ...



... and run specialized !

Projects / Collaborations

- **Delite (Stanford)** Onward!'10, PACT'11, DSL'11, IEEE Micro 10/11, ECOOP'13, GPCE'13, TECS 4/14
 - DSLs and Big Data on heterogeneous devices
- **Spiral (ETH)** GPCE'13, ARRAY'14
 - Fast numeric libraries
- **LegoBase (EPFL DATA)** DCP'14, VLDB'14
 - Databases and query processing
- **Lancet (Oracle Labs)** PLDI'14
 - Integrate LMS with JVM / JIT compilation
- **Hardware (EPFL PAL)** FPT'13, FPL'14
 - Domain specific HW synthesis
- **Parser Combinators (EPFL LAMP)** OOPSLA'14
 - Protocols and dynamic programming
- **JavaScript (EPFL, INRIA Rennes)** ECOOP'12, GPCE'13
 - LMS for the web



ORACLE®



ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

LMS = Lightweight Modular Staging

- Int, String, T
 - "execute now"
- Rep[Int], Rep[String], Rep[T]
 - "generate code to execute later"
- if (c) a else b → __ifThenElse(c,a,b)
 - "language virtualization"
- Extensible IR, transformers, loop fusion, ...
- "Batteries included"

Example: Matrix

```
class Matrix[T:Numeric:Manifest](val rows: Rep[Int], val cols: Rep[Int]) {  
    private val arr: Rep[Array[T]] = ArrayNew[T](rows * cols)  
    private val num = implicitly[Numeric[T]]; import num._  
  
    def apply(i: Rep[Int], j: Rep[Int]): A =  
        arr(i*cols + j)  
  
    def update(i: Rep[Int], j: Rep[Int], e: Rep[A]): Unit =  
        arr(i*cols + j) = e  
  
    def *(that: Matrix[T]) = {  
        val res = new Matrix[T](this.rows, that.cols)  
  
        for (i <- 0 until this.rows) {  
            for (j <- 0 until that.cols) {  
                for (k <- 0 until this.rows)  
                    res(i, j) += this(i, k) * that(k, j)  
            }  
        }  
        res  
    }  
}
```

Matrices are “now” objects, their data arrays are “later” objects

Generate Low-Level Code

```
var x27 = 500 * 500
var x28 = new Array[Double](x27)
var x29: Int = 0
while (x29 < 500) {
    var x30: Int = 0
    while (x30 < 500) {
        var x31: Int = 0
        while (x31 < 100) {
            ...
            x31 += 1
        }
        var x46 = ()
        x46
        x30 += 1
    }
    var x47 = ()
    x47
    x29 += 1
}
```

(still far from optimal:
should block loops for locality)

```
val m = randomMatrix(500, 100
val n = randomMatrix(100, 500)

val p = m * n

--- generic took 2.691s
--- generic took 1.4s
--- generic took 1.464s
--- generic took 1.359s
--- generic took 1.244s

--- double took 1.062s
--- double took 1.228s
--- double took 1.076s
--- double took 1.03s
--- double took 1.076s

--- staged took 0.088s
--- staged took 0.058s
--- staged took 0.055s
--- staged took 0.054s
--- staged took 0.056s
```

20x!

User code

```
println(...)  
val mystring = ... // Rep[String]  
println(mystring.length)
```

Sym(32) = Reflect(PrintIn(...), ...)

DSL interface

```
type Rep[T]  
def infix_length(s: Rep[String]): Rep[Int]  
def println(x: Rep[Any]): Rep[Unit]
```

Sym(45) = // mystring

Sym(46) = StrLength(Sym(45))

DSL Implementation

```
type Rep[T] = Exp[T] // Sym[T] | Const[T]
```

Sym(47) = Reflect(PrintIn(Sym(46)), List(Sym(32)))


```
case class StrLength(s: Exp[String]) extends Def[Int]  
case class PrintIn(s: Exp[Any]) extends Def[Unit]
```

```
def infix_length(s: Exp[String]): Exp[Int] = s match {  
    case Const(s) => Const(s.length)  
    case _ => reflectPure(StrLength(s)) }  
def println(x: Exp[Any]): Exp[Unit] = reflectEffect(PrintIn(x))
```

Demo Time

<https://scala-lms.github.io/tutorials/shonan.html>

<http://scala-lms.github.io/tutorials/linq.html>

A photograph of the Oracle headquarters in Redwood City, California. The image shows several modern, curved glass buildings with white structural elements. One building has "ORACLE" written on its side. The buildings are situated along a large body of water, with a rocky shoreline in the foreground. A fountain is visible in the water. The sky is clear and blue.

LMS for Efficient Data Processing

**Yannis Klonatos, Christoph Koch (EPFL),
Hassan Chafi, Tiark Rompf (Oracle)**

Databases: State of the Art

- Popular generic DBMSs consist of > 1M lines of optimized C code
- Manual specialization for performance – e.g. PostgreSQL:
 - 20 implementations of memory page abstraction
 - 7 implementations of B-trees
- Difficult to adapt
 - e.g. disk based → in memory
- Still 10 – 100x slower than hand-written queries
(Stonebraker: time for a complete rewrite, Zukowski: Monetdb/x100)
- Commercial DBMS interpret query execution plans
 - some research on query compilation using LLVM (e.g. HyPer)

LegoBase

- New in-memory DB query engine, written in Scala
- Staged query interpreter
 - Compiles query execution plans (from Oracle DB) to C code
 - Supports all 22 TPCH queries
 - ~3000 lines of Scala code
- Use LMS for additional optimizations
 - Operator inlining
 - Optimizing data structures
 - Optimizing control flow (push vs pull)

It is indeed possible to write high performance systems in a high level language



Very Large Data Base Endowment Inc.
(VLDB Endowment)

VLDB Best Paper Award

presented to

Tiark Rompf

for the paper entitled
Building Efficient Query Engines in
a High-Level Language

40th International Conference on Very Large Data Bases
September 1st-5th, Hangzhou, China

A SQL engine in 500 LOC

<https://scala-lms.github.io/tutorials/query.html>

Data is not only stored
but also transferred

Efficient, hand-optimized HTTP parser

```
switch (s) {
    case s_req_spaces_before_url:
        if (ch == '/' || ch == '*') {
            return s_req_path;
        }
        if (IS_ALPHA(ch)) {
            return s_req_schema;
        }
        break;

    case s_req_schema:
        if (IS_ALPHA(ch)) {
            return s;
        }

        if (ch == ':') {
            return s_req_schema_slash;
        }
        break;
```

- Originally part of Nginx, later Node.js
- 2000+ lines of code
- Callbacks for header names/values triggered
- State-machine like code
- “Flat” code, loops/conditions

Staged Parser Combinators

```
def status: Parser[Int] =  
  ("HTTP/"~decimalNumber)~>wholeNumber<~(wildRegex~crlf) ^^ (_.toInt)  
  
def header: Parser[Option[(String,Any)]] =  
  (headerName~~":")~(wildRegex~crlf) ^^ {  
    case hName~prop => collect(hName.toLowerCase, prop)  
  }  
  
def headers = rep(header)  
  
def response = status ~ headers  
  
...
```

- 200+ lines of code
- Fairly easy to change behaviour of a parser
 - Ex: \sim vs $\langle \sim$ vs $\sim \rangle$

Staged Parser Combinators

